Basic Approaches for Missing Values

Basic Approaches for Missing Values

We will now discuss a few simple approaches to deal with missing values

We will use partially synthetic data

- We will focus on specific (and mosly intact) sections of our series

- The we will remove values artificially

- ...And measure the accuracy of our filling approaches via the Root MSE

$$RMSE = \sqrt{\frac{1}{n} \sum_{i=0}^{n} (x_i - \hat{x}_i)^2}$$

Where x_i is a value from the filled series and $\hat{x}_i$ the ground truth

- $x_i = \hat{x}_i$ if no value is missing

- Hence, any MSE difference is entirely due to missing values

The Benchmark Dataset

Our benchmark dataset will consist of this particular stretch from our traffic series

```
In [2]:  util.plot_series(ddata, figsize=figsize)
```


- There are (comparatively) few missing values
- Some of them are isolated, some for contiguous "holes"

The Benchmark Dataset

We now introduce some missing values

…By drawing them at random:

```python
In [3]: np.random.seed(42) # seed (to get reproducible results)
mv_idx = np.random.choice(range(len(ddata.index)), size=30, replace=False)
ddata_mv = ddata.copy()
ddata_mv.iloc[mv_idx] = np.NaN
util.plot_series(ddata_mv, figsize=figsize)
plt.scatter(ddata.index[mv_idx], ddata.iloc[mv_idx], color='tab:orange', marker='x');
```

Forward/Backward Filling

The easiest approach for missing values consists in replicating nearby observations

- Forward filling: propagate forward the last valid observation

- Backward filling: propagate backward the next valid observation

An important observation:

- When filling missing values, we have access to the whole series

- ...So we can reason both forward and backward

Forward/backward filling are simple methods, but they can work well

- Rationale: most time series have a certain "inertia"

- ...I.e.: a strong level of local correlation

- For this reason (e.g.) the last observation is often a good predictor for the next one

Forward/Backward Filling

Forward and backward filling are pre-implemented in pandas

They are available through the `fillna` method:

```
DataFrame.fillna(..., method=None, ...)
```

- `fillna` replaces `NaN` values in a `DataFrame` or `Series`

- The `method` parameter can take the values:

  - "pad" or "ffill": these correspond to forward filling

  - "backfill" or "bfill": these correspond to backward filling

They are generally applied to datasets with a dense index

- Remember that our benchmark dataset already has a dense index

Forward/Backward Filling on the Benchmark

We can finally test forward/backward filling

```
In [4]:  nan_mask = ddata['value'].isnull()
         ffseries = ddata_mv.fillna(method='ffill')
         ffseries[nan_mask] = np.NaN # We empty the values that were originally empty
         bfseries = ddata_mv.fillna(method='bfill')
         bfseries[nan_mask] = np.NaN # We empty the values that were originally empty
```

We can check the corresponding RMSE:

```
In [6]:  rmse_ff = np.sqrt(mean_squared_error(ddata[~nan_mask], ffseries[~nan_mask]))
         rmse_bf = np.sqrt(mean_squared_error(ddata[~nan_mask], bfseries[~nan_mask]))
         print(f'RMSE for forwad filling: {rmse_ff:.2f}, for backward filling {rmse_bf:.2f}')

         RMSE for forwad filling: 1.33, for backward filling 0.87
```

- In this case backward filling seems to work better

- The results are of course application-dependent

Forward/Backward Filling on the Benchmark

Let's have a close look at the results for forward filling

```
In [16]:  util.plot_series(ffseries, figsize=figsize)
          plt.scatter(ddata.index[mv_idx], ddata.iloc[mv_idx], color='tab:orange', marker='x');
          plt.scatter(ddata.index[mv_idx], ffseries.iloc[mv_idx], color='tab:red', marker='x');
```


- The red marks represent the values that have been filled

Forward/Backward Filling on the Benchmark

Let's have a close look at the results for backward filling

```
In [17]:  util.plot_series(bfseries, figsize=figsize)
          plt.scatter(ddata.index[mv_idx], ddata.iloc[mv_idx], color='tab:orange', marker='x');
          plt.scatter(ddata.index[mv_idx], bfseries.iloc[mv_idx], color='tab:red', marker='x');
```


- Forward/backward filling tend to work well for low variance sections
- ...And conversely work worse for high variance sections

(Geometric) Interpolation

A few more options are available via the `interpolate` method

```
DataFrame/Series.interpolate(method='linear', ...)
```

The `method` parameter determines how NaNs are filled:

- "linear" uses a linear interpolation, assuming uniformly spaced samples
- "time" uses a linear interpolation, but supports non-uniformly spaced samples
- "nearest" uses the closest value
- "polynomial" uses a polynomial interpolation
- Even "ffill" and "bfill" are available

Both "polynomial" and "spline" require to specify the additional parameter `order`

- E.g. `df.interpolate(method='polynomial', order='3')`

(Geometric) Interpolation

Let us check the performance of some approaches

```python
In [45]: args = [{'method':'linear'}, {'method':'time'}, {'method':'nearest'},
              {'method':'polynomial', 'order': 2}, {'method':'spline', 'order': 4}]
         filling_res = {}
         for a in args:
             filling_res[a['method']] = ddata_mv.interpolate(**a)
             rmse = np.sqrt(mean_squared_error(ddata[~nan_mask], filling_res[a['method']][~nan_mask]))
             print(f'RMSE for {a["method"]}: {rmse:.2f}')

RMSE for linear: 0.98
RMSE for time: 0.98
RMSE for nearest: 1.35
RMSE for polynomial: 1.10
RMSE for spline: 1.03
```

- "linear" and "time" are equivalent (we have uniformly-spaced samples)

- "polynomial" is the most complex, and in this case also the worst

All perform worse than backward filling (at least in this case)!

Linear Filling on the Benchmark

Let's have a close look at the results for linear filling

```
In [46]: util.plot_series(filling_res['linear'], figsize=figsize)
         plt.scatter(ddata.index[mv_idx], ddata.iloc[mv_idx], color='tab:orange', marker='x');
         plt.scatter(ddata.index[mv_idx], filling_res['linear'].iloc[mv_idx], color='tab:red', marker='x'
```


- Linear filling works well for series with slower dynamics
- ...But does not work well for series that faster dynamics

Nearest Filling on the Benchmark

Let's have a close look at the results for nearest filling

```
In [47]: util.plot_series(filling_res['nearest'], figsize=figsize)
         plt.scatter(ddata.index[mv_idx], ddata.iloc[mv_idx], color='tab:orange', marker='x');
         plt.scatter(ddata.index[mv_idx], filling_res['nearest'].iloc[mv_idx], color='tab:red', marker='>
```


- Nearest filling is a compromise between forward and backward filling
- ...And in our case its performance is the average of the two

Polynomial Filling on the Benchmark

Let's have a close look at the results for polynomial filling

```python
util.plot_series(filling_res['polynomial'], figsize=figsize)
plt.scatter(ddata.index[mv_idx], ddata.iloc[mv_idx], color='tab:orange', marker='x');
plt.scatter(ddata.index[mv_idx], filling_res['polynomial'].iloc[mv_idx], color='tab:red', marker
```


- Polynomial interpolation relies in nearby value to fit a polynomial
- High-order polinomial often vary too much and work less well

Spline Filling on the Benchmark

Let's have a close look at the results for spline filling

```
In [49]: util.plot_series(filling_res['spline'], figsize=figsize)
         plt.scatter(ddata.index[mv_idx], ddata.iloc[mv_idx], color='tab:orange', marker='x');
         plt.scatter(ddata.index[mv_idx], filling_res['spline'].iloc[mv_idx], color='tab:red', marker='x'
```


- Spline interpolation relies on piecewise polynomial curves

- ...And it's often more robust than polynomial interpolation

How do we choose which method is best? Is the (R)MSE enough?